

# Building a Kubernetes Cluster

## Relevant Documentation

- [Installing kubeadm](#)
- [Creating a cluster with kubeadm](#)

## Lesson Reference

If you are using cloud playground, create three servers with the following settings:

- Distribution: Ubuntu 18.04 Bionic Beaver LTS
- Size: medium

On all nodes, set up containerd. You will need to load some kernel modules and modify some system settings as part of this process.

```
cat <<EOF | sudo tee /etc/modules-load.d/containerd.conf
overlay
br_netfilter
EOF

sudo modprobe overlay

sudo modprobe br_netfilter

cat <<EOF | sudo tee /etc/sysctl.d/99-kubernetes-cri.conf
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
net.bridge.bridge-nf-call-ip6tables = 1
EOF

sudo sysctl --system
```

Install and configure containerd.

```
sudo apt-get update && sudo apt-get install -y containerd

sudo mkdir -p /etc/containerd

sudo containerd config default | sudo tee /etc/containerd/config.toml

sudo systemctl restart containerd
```

On all nodes, disable swap.

```
sudo swapoff -a

sudo sed -i '/ swap / s/^(.*)$/#\1/g' /etc/fstab
```

On all nodes, install kubeadm, kubelet, and kubectl.

```
sudo apt-get update && sudo apt-get install -y apt-transport-https curl

curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -

cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF

sudo apt-get update

sudo apt-get install -y kubelet=1.20.1-00 kubeadm=1.20.1-00 kubectl=1.20.1-00

sudo apt-mark hold kubelet kubeadm kubectl
```

**On the control plane node only**, initialize the cluster and set up kubectl access.

```
sudo kubeadm init --pod-network-cidr 192.168.0.0/16

mkdir -p $HOME/.kube

sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Verify the cluster is working.

```
kubectl version
```

Install the Calico network add-on.

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
```

Check the calico-related `kube-system` Pods to verify that everything is working so far (they may take a few moments to fully start up).

```
kubectl get pods -n kube-system
```

Get the join command (this command is also printed during `kubeadm init` . Feel free to simply copy it from there).

```
kubeadm token create --print-join-command
```

Copy the join command from the control plane node. Run it on each worker node as root (i.e. with `sudo` ).

```
sudo kubeadm join ...
```

On the control plane node, verify all nodes in your cluster are ready. Note that it may take a few moments for all of the nodes to enter the `READY` state.

```
kubectl get nodes
```